

Finite differences as approximations of partial derivatives




Figure: Select every all rows of u and from the first column to the second.




Figure: Select from the second row to the second and from the first column to the second.

Estimate the error for $\partial^+ u(x)$

By Taylor's expansion, for some $\xi \in [x, x+h]$

$$(4) \quad \partial_x u(x) = \partial^+ u(x) - \frac{\Delta x}{2} \partial_x^2 u(\xi).$$




Figure: Error of $\partial^+ u(x)$ for several values of Δx 🍷.

Part of the error is due to the inaccuracies in (4) is

$$\text{truncated error} = \frac{\Delta x}{2} \left| \partial_x^2 u(\xi) \right|.$$

Depends of $\varepsilon_{\text{mach}}$ (for float64 is $2.22044604925 \times 10^{-16}$).

$$\text{rounded error} \approx \frac{|u(x)| \varepsilon_{\text{mach}}}{\Delta x} + \left| \partial_x u(x) \right| \varepsilon_{\text{mach}}.$$

The best value of Δx is obtained by minimizing the total error as function of Δx , i.e.,

$$0 = \frac{1}{2} \left| \partial_x^2 u(\xi) \right| - \frac{|u(x)| \varepsilon_{\text{mach}}}{(\Delta x)^2}.$$
$$\Delta x = \sqrt{\frac{2 |u(x)| \varepsilon_{\text{mach}}}{\left| \partial_x^2 u(\xi) \right|}}.$$

If $u(x)$ and $\partial_x^2 u(\xi)$ are neither large nor small, then

$$\Delta x \approx \sqrt{\varepsilon_{\text{mach}}} \quad (\text{for float64 is } 1.49011611938 \times 10^{-8}).$$

Truncation error

$$u(x + \Delta x) = \underbrace{u(x) + \Delta x \partial_x u(x)}_{1^{\text{st}} \text{ order approximation}} + \underbrace{\frac{(\Delta x)^2}{2!} \partial_x^2 u(x) + \dots}_{\text{truncation error}}.$$

$$u(x + \Delta x) = \text{approximation} + E(\Delta x) + \text{high order terms.}$$

$$u(x + \Delta x) \approx u(x) + \Delta x \partial_x u.$$

$$E(\Delta x) = \frac{(\Delta x)^{n+1}}{(n+1)!} \partial_x^{n+1} u(x).$$

```
import numpy as np
```

```
x = 0
```

```
Δx = 0.1
```

```
u = np.cos
```

```
def dudx(x):
    return -np.sin(x)
```

```
def du2dx(x):
    return -np.cos(x)
```

```
exact = u(x + Δx)
approximation = u(x) + Δx * dudx(x)
absolute_lower_order_omitted_term = np.abs(np.power(Δx, 2) / 2 * du2dx(x))
truncation_error = np.abs(exact - approximation)
error_attributed_to_other_omitted_terms = np.abs(
    truncation_error - absolute_lower_order_omitted_term
)
```

```
First order approximation: 1.0
```

```
Exact: 0.9950041652780258
```

```
Truncation error: 0.0049958347219741794
```

```
Absolute lower order omitted term: 0.0050000000000000001
```

```
Error attributed to other terms: 4.165278025821534e-06
```

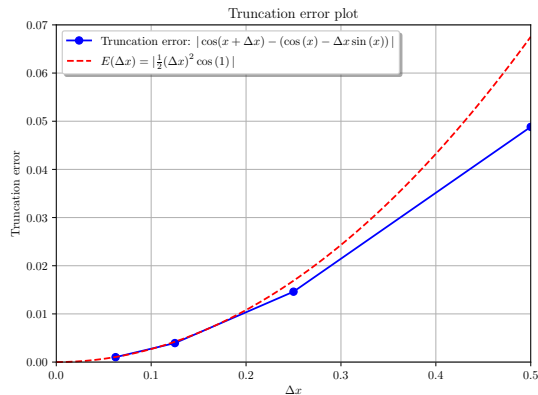


Figure: Truncation error for several values of Δx at $x = 1$.

Finite differences as approximations of partial derivatives

Given a known truncation error $E(\Delta x) = O((\Delta x)^n)$ it is possible to estimate the **order** of an approximation.

$$E(\Delta x) \approx C(\Delta x)^n.$$

Taking logarithms of both sides gives

$$\log |E(\Delta x)| = n \log(\Delta x) + \log(C).$$

which is a linear function in Δx . Therefore, one way of approximate the slope n as a gradient is

$$n \approx \frac{\log |E(\max(\Delta x))| - \log |E(\min(\Delta x))|}{\log(\max(\Delta x)) - \log(\min(\Delta x))}.$$

Δx	backward	centered	forward
0.1	-0.670603	-0.705929	-0.741255
0.05	-0.689138	-0.706812	-0.724486
0.025	-0.698195	-0.707033	-0.715872
0.0125	-0.702669	-0.707088	-0.711508

Δx	backward	centered	forward
0.1	0.0365038	0.00117792	0.034148
0.05	0.0179686	0.000294591	0.0173794
0.025	0.00891203	7.36547e-05	0.00876472
0.0125	0.00443777	1.84141e-05	0.00440095

Order of convergence of backward is 1.013379633444132
Order of convergence of centered is 1.9997633049971728
Order of convergence of forward is 0.9853046896368071

```
import numpy as np
from jaxtyping import Array, Float

u = np.cos

def dudx(x):
    return -np.sin(x)

x = np.pi / 4
Δx = np.logspace(start=-3, stop=0, num=4, base=2) / 10

backward: Float[Array, "dim1"] = (u(x) - u(x - Δx)) / Δx
centered: Float[Array, "dim1"] = (u(x + Δx) - u(x - Δx)) / (2 * Δx)
forward: Float[Array, "dim1"] = (u(x + Δx) - u(x)) / Δx



error_backward: Float[Array, "dim1"] = np.abs(dudx(x) - backward)
error_centered: Float[Array, "dim1"] = np.abs(dudx(x) - centered)
error_forward = np.abs(dudx(x) - forward)

def estimate_order(
    Δx: Float[Array, "dim1"], truncation_error: Float[Array, "dim1"]
) → float:
    assert Δx.size == truncation_error.size

    E_max = truncation_error[np.argmax(a=Δx)]
    E_min = truncation_error[np.argmin(a=Δx)]

    return (np.log(np.abs(E_max)) - np.log(np.abs(E_min))) / (
        np.log(Δx.max()) - np.log(Δx.min())
    )

print(f"Order of convergence of backward is {estimate_order(Δx, error_backward)}")
print(f"Order of convergence of centered is {estimate_order(Δx, error_centered)}")
print(f"Order of convergence of forward is {estimate_order(Δx, error_forward)}")
```

Program  : Determine order of convergence .

Order of convergence

```
from typing import Callable

import numpy as np

def u(x: float) → float:
    """Sample function
     $u: \mathbb{R} \rightarrow \mathbb{R}$ 
     $x \mapsto \exp(x^2)$ 
    """
    return np.exp(np.pow(x, 2))

def up(x: float) → float:
    """Derivative of sample function
     $u': \mathbb{R} \rightarrow \mathbb{R}$ 
     $x \mapsto 2*x*f(x)$ 
    """
    return 2 * x * u(x)

def ff(u: Callable, Δx: np.array) → np.array:
    """Forward finite difference approximation
     $u' = (u(x + \Delta x) - u(x)) / \Delta x$ 
    """
    return (u(x + Δx) - u(x)) / Δx

def bf(u: Callable, Δx: np.array) → np.array:
    """Backward finite difference approximation
     $u' = (u(x) - u(x - \Delta x)) / \Delta x$ 
    """
    return (u(x) - u(x - Δx)) / Δx

def cf(u: Callable, Δx: np.array) → np.array:
    """Centered finite difference approximation
     $u' = (u(x + \Delta x / 2) - u(x - \Delta x / 2)) / \Delta x$ 
    """
    return (u(x + Δx / 2) - u(x - Δx / 2)) / Δx

x = 2
exact = up(x)
Δx = np.logspace(start=-16, stop=-1.0, num=16)
```

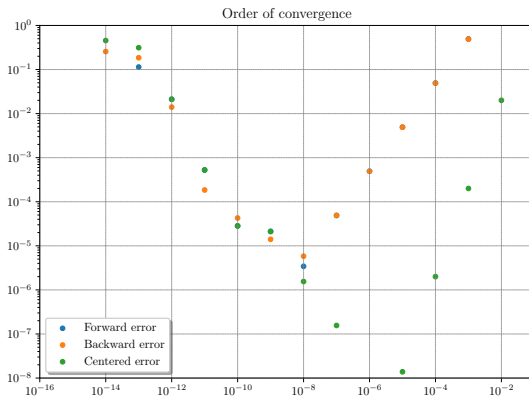


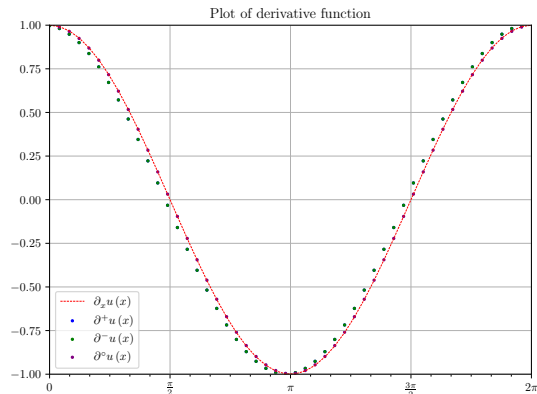
Figure: $\partial^+ u(x)$, $\partial^- u(x)$, $\partial^\circ u(x)$ of $u(x) = \exp(x^2)$ for several values of Δx at $x = 2$ \oplus .

Finite differences as approximations of partial derivatives

In the array computation, we do not consider the last element of the difference $u_{i+1} - u_i$. and the first element of the difference $u_{i+1} - u_i$.

$$\partial^+ u(x) = \frac{u_{i+1} - u_i}{\Delta x}, \quad i = 0, \dots, n-1.$$

$$\partial^- u(x) = \frac{u_i - u_{i-1}}{\Delta x}, \quad i = 1, \dots, n.$$



```
import numpy as np
```

```
x, Δx = np.linspace(start=0, stop=2 * np.pi, retstep=True)
y = np.sin(x)
```

```
forward = (np.roll(y, -1) - y)[: -1] / Δx
```

```
backward = (y - np.roll(y, 1))[1:] / Δx
```

```
centered = (np.roll(y, -1) - np.roll(y, 1))[1: -1] / (2 * Δx)
```

```
first_derivative = np.cos(x)
```

```
np.roll(a = u, shift = -1), u, np.roll(a = u, shift = -1) - u
```

u[1]	u[2]	u[3]	u[4]	u[0]
u[0]	u[1]	u[2]	u[3]	u[4]
u[1] - u[0]	u[2] - u[1]	u[3] - u[2]	u[4] - u[3]	u[0] - u[4]

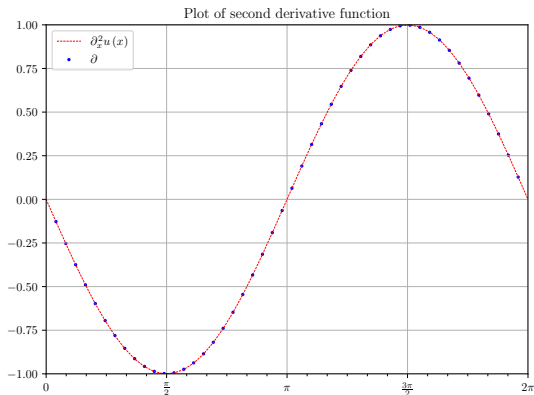
```
np.roll(a = u, shift = 1), u, u - np.roll(a = u, shift = 1)
```

u[4]	u[0]	u[1]	u[2]	u[3]
u[0]	u[1]	u[2]	u[3]	u[4]
u[4] - u[0]	u[0] - u[1]	u[1] - u[2]	u[2] - u[3]	u[3] - u[4]

Finite differences as approximations of partial derivatives

$$\partial^+ \partial^- u(x) = \frac{u_{i+1} - 2u_i + u_{i-1}}{(\Delta x)^2}, \quad i = 1, \dots, n-1.$$

```
import numpy as np
x, Δx = np.linspace(start=0, stop=2 * np.pi, retstep=True)
y = np.sin(x)
plt.ylim(-1, 1)
plt.grid()
```




Finite differences as approximations of partial derivatives

$$= \frac{u_{i+2} - 2u_{i+1} + 2u_{i-1} - u_{i-2}}{2(\Delta x)^3}, \quad i = 2, \dots, n-2.$$

```
import numpy as np
```

```
x, Δx = np.linspace(start=0, stop=2 * np.pi, retstep=True)
y = np.sin(x)
```

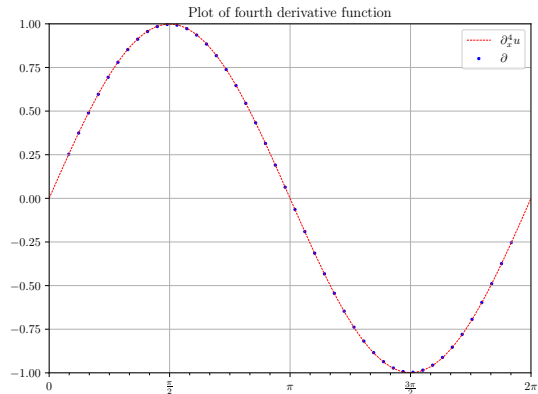
```
plt.gca().xaxis.set_minor_locator(plt.MultipleLocator(np.pi / 12))
plt.gca().xaxis.set_major_formatter(plt.FuncFormatter(multiple_formatter)))
plt.scatter(x=x[1:-1], y=partial2x, c="blue", s=3, label=r"$\partial^2 u$")
plt.title(label="Plot of second derivative function")
```



Finite differences as approximations of partial derivatives

$$= \frac{u_{i+2} - 4u_{i+1} + 6u_i - 4u_{i-1} + u_{i-2}}{(\Delta x)^4}, \quad i = 2, \dots, n-2.$$

```
import numpy as np
x, Δx = np.linspace(start=0, stop=2 * np.pi, retstep=True)
y = np.sin(x)
plt.gca().xaxis.set_minor_locator(plt.MultipleLocator(np.pi / 12))
plt.gca().xaxis.set_major_formatter(plt.FuncFormatter(multiple_formatter)))
plt.scatter(x=x[2:-2], y=partial3x, c="blue", s=3, label=r"$\partial^3$")
plt.title(label="Plot of third derivative function")
```



Definition (Holomorphic function)

Let $D \subset \mathbb{C}$ be a simply connected, open region and $u: D \rightarrow \mathbb{C}$. We say that u is **complex differentiable** at $a \in D$ iff

$$\lim_{z \rightarrow a} \frac{u(z) - u(a)}{z - a}$$

exists. If u is complex differentiable at every point of D , then we say that u is **holomorphic** in D .

The **complex step derivative** approximation is a technique to compute the derivative of a real-valued function $u(x)$. For u analytic,

$$u(x + i\Delta x) = u(x) + i\Delta x \partial_x u(x) + \frac{(i\Delta x)^2}{2!} \partial_x^2 u(x) + \dots$$

$$\operatorname{Re}[u(x + i\Delta x)] + i \operatorname{Im}[u(x + i\Delta x)] \approx u(x) + i\Delta x \partial_x u(x).$$

Comparing imaginary parts of the two sides gives

$$\partial_x u(x) \approx \operatorname{Im} \left[\frac{u(x + i\Delta x)}{\Delta x} \right].$$

Remark

Behind the scenes, the complex step method is a particular case of **automatic differentiation**.

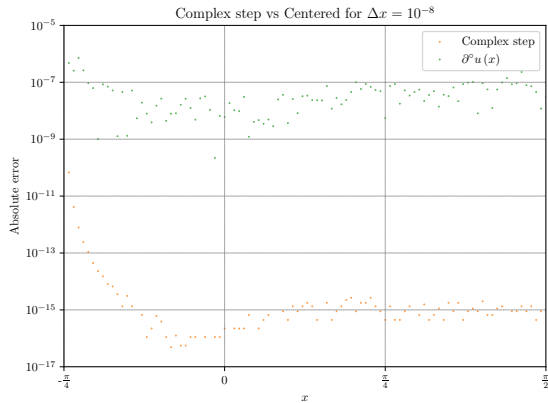

$$u(x + i\Delta x) = u(x) + i\Delta x \partial_x u(x) + \frac{(i\Delta x)^2}{2!} \partial_x^2 u(x) + \cdots.$$

$$\operatorname{Re}[u(x + i\Delta x)] + i \operatorname{Im}[u(x + i\Delta x)] \approx u(x) + i\Delta x \partial_x u(x).$$

Comparing real parts of the two sides gives


$$\partial_x^2 u(x) \approx \frac{2}{\Delta x^2} (u(x) - \operatorname{Re}[u(x + i\Delta x)]).$$

Complex step method



Complex step method

Complex step vs Centered with varying step size



Solve BVP for ODEs

Let's follow the steps:

- 1 Discretize the domain on which the equation is defined.
- 2 On each **grid point**, replace the derivatives with an approximation, using the values in neighbouring grid points.
- 3 Replace the exact solutions by their approximations.
- 4 Solve the resulting system of equations.

Finite difference for Two-point BVP

We will first see how to find approximations to the derivative of a function, and then how these can be used to solve boundary value problems like

$$\begin{cases} \frac{d^2 u}{dx^2} + p(x) \frac{du}{dx} + q(x) u = r(x) & \text{for } a \leq x \leq b. \\ u(a) = u_a, \quad u(b) = u_b \end{cases}.$$

This technique described here is applicable to several other time dependent PDEs, and it is therefore important to try to understand the underlying idea.

Example (Two-point BVP FDM for the 1D Poisson Problem)

Let $f: [0, 1] \rightarrow \mathbb{R}$ be a function. Find a $u: [0, 1] \rightarrow \mathbb{R}$ such that

$$(5) \quad \begin{cases} -\frac{d^2 u}{dx^2} = f(x), & x \in (0, 1). \\ u(0) = u_a, \quad u(1) = u_b. \end{cases}$$

Instead of trying to compute $u(x)$ exactly, we will now try to compute a numerical approximation u_Δ of $u(x)$. As many times before we start by defining $n+1$ equally spaced points $\{x_i\}_{i=0}^n$ with a grid size $h = \frac{b-a}{n}$ so that

$$\forall i = 0, 1, \dots, n : x_i := a + ih.$$

Consider a collection of equally spaced points, labeled with an index i , with the physical spacing between them denoted Δx . We can express the first derivative of a quantity a at i as:

$$\frac{\partial a}{\partial x_i} \approx \frac{a_i - a_{i-1}}{\Delta x}$$


or

$$\frac{\partial a}{\partial x_i} \approx \frac{a_{i+1} - a_i}{\Delta x}$$

$$a_{i+1} = a_i + \Delta x \frac{\partial a}{\partial x} \Big|_i + \frac{1}{2} \Delta x^2 \frac{\partial^2 a}{\partial x^2} \Big|_i + \dots$$

Solving for $\partial a / \partial x|_i$, we see

$$\begin{aligned} \frac{\partial a}{\partial x} \Big|_i &= \frac{a_i - a_{i-1}}{\Delta x} - \frac{1}{2} \Delta x \frac{\partial^2 a}{\partial x^2} \Big|_i + \dots \\ &= \frac{a_i - a_{i-1}}{\Delta x} + \mathcal{O}(\Delta x) \end{aligned}$$




```
import numpy as np

from scipy.sparse import csr_array, diags_array
from scipy.sparse.linalg import spsolve

def fdm_poisson1d_matrix(N: int):
    """Computes the finite difference matrix for the Poisson problem in 1D

    Parameters:
    N (int): Number of grid points :math:\\{x_i\\}_{i=0}^N` counting from 0.

    Returns:
    A (scipy.sparse.csr.csr_array): Finite difference sparse matrix

    """
    Δx = 1 / N
    diag = np.concatenate(
        (
            np.ones(shape=1),
            np.full(shape=N - 1, fill_value=2 / Δx**2),
            np.ones(shape=1),
        )
    )
    diag_sup = np.concatenate(
        (np.zeros(shape=1), np.full(shape=N - 1, fill_value=-1 / Δx**2))
    )
    diag_inf = np.flipud(m=diag_sup)

    return diags_array(
        [diag, diag_sup, diag_inf],
        offsets=[0, 1, -1],
        shape=(N + 1, N + 1),
        format="csr",
    )


N = 10
x = np.linspace(start=0, stop=1, num=N + 1)
A = fdm_poisson1d_matrix(N)
F = (2 * np.pi) ** 2 * np.sin(2 * np.pi * x)
```

Program : fdmpoisson1d.py.

```
xfine = np.linspace(start=0, stop=1, num=10 * N)
# Analytical reference solution
u = np.sin(2 * np.pi * xfine)

# Incorporate boundary condition into rhs vector
F[0], F[-1] = u[0], u[-1]

# Solve AU = F
U = spsolve(A=A, b=F)
```



```
%%MatrixMarket matrix coordinate real general
%
```

```
11 11 29
1 1 1
2 1 -9.999999999999999E1
2 2 1.9999999999999997E2
2 3 -9.999999999999999E1
3 2 -9.999999999999999E1
3 3 1.9999999999999997E2
3 4 -9.999999999999999E1
4 3 -9.999999999999999E1
4 4 1.9999999999999997E2
4 5 -9.999999999999999E1
5 4 -9.999999999999999E1
5 5 1.9999999999999997E2
5 6 -9.999999999999999E1
6 5 -9.999999999999999E1
6 6 1.9999999999999997E2
6 7 -9.999999999999999E1
7 6 -9.999999999999999E1
7 7 1.9999999999999997E2
7 8 -9.999999999999999E1
8 7 -9.999999999999999E1
8 8 1.9999999999999997E2
8 9 -9.999999999999999E1
9 8 -9.999999999999999E1
9 9 1.9999999999999997E2
9 10 -9.999999999999999E1
10 9 -9.999999999999999E1
10 10 1.9999999999999997E2
10 11 -9.999999999999999E1
11 11 1
```

Program  : poissonA.mm.

```
%%MatrixMarket matrix coordinate real general
%
```

```
1 11 10
1 2 2.3204831651684845E1
1 3 3.754620631564544E1
1 4 3.754620631564544E1
1 5 2.3204831651684845E1
1 6 4.8347117754578846E-15
1 7 -2.3204831651684856E1
1 8 -3.754620631564544E1
1 9 -3.754620631564544E1
1 10 -2.3204831651684856E1
1 11 -2.4492935982947064E-16
```

Program  : poissonF.mm.

```
%%MatrixMarket matrix coordinate real general
%
```

```
1 11 10
1 2 6.07510379673303E-1
1 3 9.829724428297576E-1
1 4 9.829724428297576E-1
1 5 6.07510379673303E-1
1 6 -5.1532467934608046E-17
1 7 -6.075103796733031E-1
1 8 -9.829724428297577E-1
1 9 -9.829724428297578E-1
1 10 -6.075103796733033E-1
1 11 -2.4492935982947064E-16
```

Program  : poissonU.mm.




Figure: Solution.

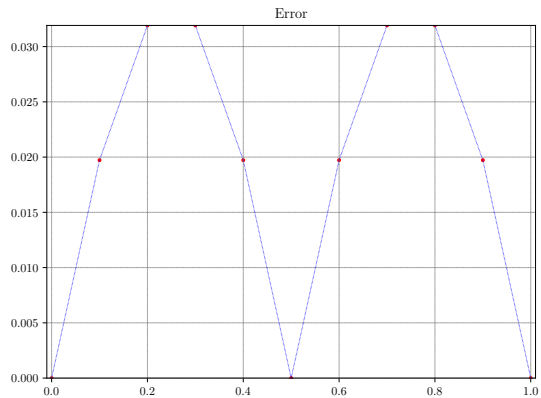


Figure: Error.

```

import numpy as np

from scipy.sparse import csr_array, diags_array
from scipy.sparse.linalg import spsolve

def tridiag(p: np.ufunc, q: np.ufunc, N: int):
    """
    Help function
    Returns a tridiagonal matrix A of dimension N+1 x N+1.
    """
    Δx = 1 / N
    diag = np.concatenate(
        (
            np.ones(shape=1),
            np.full(shape=N - 1, fill_value=-2 + Δx**2 * q),
            np.ones(shape=1),
        )
    )
    diag_sup = np.concatenate(
        (np.zeros(shape=1), np.full(shape=N - 1, fill_value=1 + Δx / 2 * p))
    )
    diag_inf = np.concatenate(
        (np.full(shape=N - 1, fill_value=1 - Δx / 2 * p), np.zeros(shape=1))
    )

    return diags_array(
        [diag, diag_sup, diag_inf],
        offsets=[0, 1, -1],
        shape=(N + 1, N + 1),
        format="csr",
    )

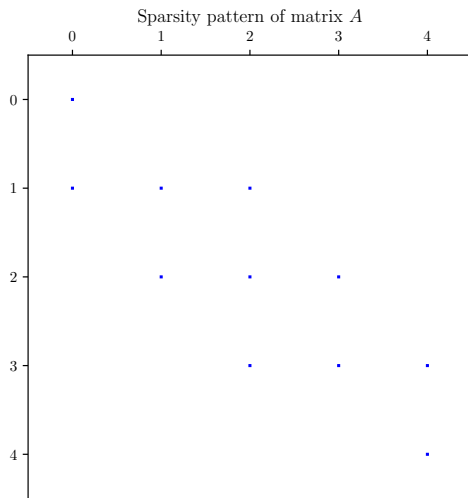
N = 4 # Number of intervals
x, Δx = np.linspace(start=0, stop=1, num=N + 1, retstep=True)

p = 2
q = -3
r = 9 * x

A = tridiag(p, q, N)
b = Δx**2 * r
b[0] = 1
b[N] = np.exp(-3) + 2 * np.exp(1) - 5

```

```
U = spsolve(A=A, b=b) # Solve the equation
```



Program : twopointboundary.py.

```
%%MatrixMarket matrix coordinate real general
```

```
%  
5 5 11  
1 1 1  
2 1 7.5E-1  
2 2 -2.1875  
2 3 1.25  
3 2 7.5E-1  
3 3 -2.1875  
3 4 1.25  
4 3 7.5E-1  
4 4 -2.1875  
4 5 1.25  
5 5 1
```

Program  : twopointboundaryA.mm.

```
%%MatrixMarket matrix coordinate real general
```

```
%  
1 5 5  
1 1 1  
1 2 1.40625E-1  
1 3 2.8125E-1  
1 4 4.21875E-1  
1 5 4.863507252859538E-1
```

Program  : twopointboundaryb.mm.

```
%%MatrixMarket matrix coordinate real general
```

```
%  
1 5 5  
1 1 1  
1 2 2.931756779400817E-1  
1 3 2.5557436395142973E-2  
1 4 9.382010692745119E-2  
1 5 4.863507252859538E-1
```

Program  : twopointboundaryU.mm.

